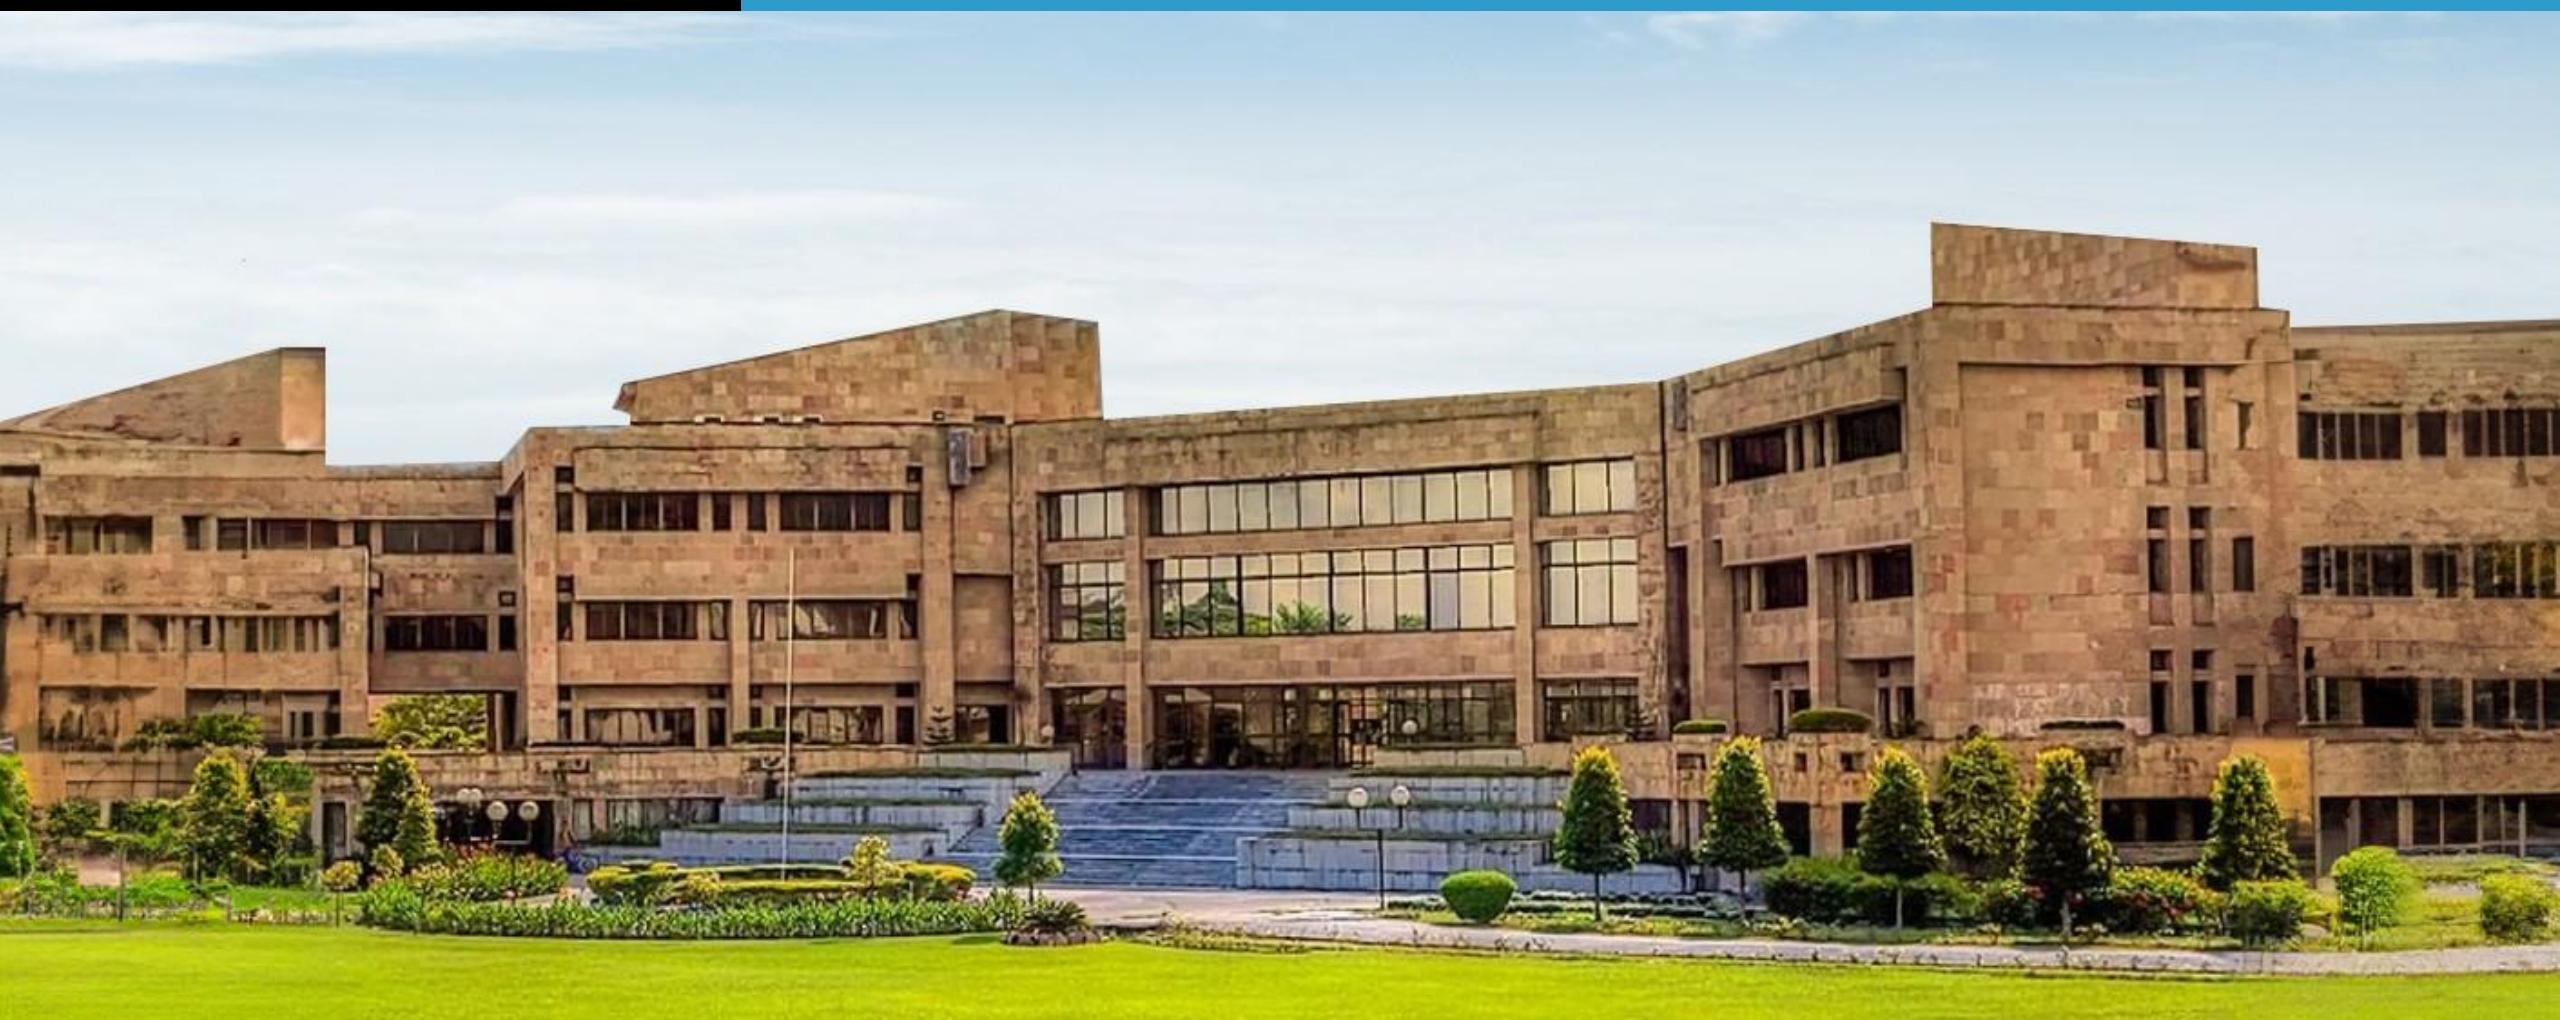


IIT Guwahati

# Advanced Certification Programme in Data Science Business Analytics



# Week 20

## Predictive Modelling



# Topics Covered

- Implementing Linear Regression
- Implementing Logistic Regression
- Q & A

# Async Recap

## 1. Understand Data Science Foundations

Master core mathematical and statistical concepts like linear algebra, probability, and descriptive statistics to support data discovery and prediction.

## 2. Apply Linear Algebra to Data

Use vectors, matrices, and transformations to handle complex data structures and optimise machine learning models.

## 3. Use Key Operations in Data Handling

Perform essential vector and matrix operations such as addition, multiplication, and transposition for efficient data analysis.

## 4. Summarise and Visualise with Statistics

Describe data using statistical summaries and visual tools like bar, pie, and box plots to identify trends and patterns.

## 5. Assess Risk Using Probability

Quantify uncertainty with probability to support decision-making and evaluate event likelihoods.

# Implementing Linear Regression

---

# Introduction to Linear Regression with Statsmodels

## Modelling Relationships and Analysing Predictors



- Model the relationship between dependent and independent variables
- Provide detailed statistical summaries
- Support hypothesis testing and confidence intervals
- Prefer for in-depth regression over Scikit-Learn
- Analyse effect of multiple predictors on house prices

# Dataset Overview

## Features Used to Predict House Price

Feature	Description
Area	Square footage of the house
Bedrooms	Number of bedrooms
Bathrooms	Number of bathrooms
Material	Type of construction material (Concrete/Masonry)
Locality	Location of the house
Price	House price (dependent variable)

- Build a linear regression model to predict price based on features

# Importing Libraries and Loading Data

## Preparing Data for Linear Regression

### Code implementation

```
import pandas as pd
import numpy as np
import statsmodels.api as sm
from sklearn.preprocessing import
OneHotEncoder

# Load dataset
data = pd.read_csv('house.csv')

# One-hot encode categorical variables
data = pd.get_dummies(data,
    columns=['Material', 'Locality'],
    drop_first=True)
```

### Selecting features and adding a constant term

```
# Selecting features and target variable
X = data[['Area', 'Bedrooms', 'Bathrooms',
    'Material_Masonry', 'Locality_Summit View']]
y = data['Price']

# Adding constant term for intercept
X = sm.add_constant(X)
```

# Building and Fitting the Model

## Creating and Training an OLS Regression Model

### Build and fit model

- Create OLS model and fit it to the data

```
# Building the model  
model = sm.OLS(y, X)
```

```
# Fitting the model  
results = model.fit()
```

# Model Summary and Interpretation

## Understanding Key Outputs from the Regression Model



- **R-squared:** Measures goodness of fit
- **F-statistic and p-value:** Determines overall model significance
- **Coefficients and standard errors:** Indicates impact and reliability of predictors
- **P-values:** Test significance of individual predictors

```
# Model summary
```

```
print(results.summary())
```

# Evaluating Model Performance

## Assessing Predictions and Error Metrics

### Make predictions

Generate predicted values using the fitted model

### Calculate residuals

Compute error between actual and predicted

### Compute error metrics

Calculate MSE and RMSE to evaluate model performance

#### # Making predictions

```
y_pred = results.predict(X)
```

#### # Calculating residuals

```
residuals = y - y_pred
```

```
from sklearn.metrics import  
mean_squared_error
```

```
mse = mean_squared_error(y,  
y_pred)
```

```
rmse = np.sqrt(mse)  
print(f'MSE: {mse}, RMSE:  
{rmse}')
```

# Checking Regression Assumptions

## Ensuring Validity of Model Results

### Linearity:

- Ensure relationship between variables is linear
- Check with scatterplot of residuals vs predicted values

### Independence:

- Ensure residuals are independent
- Use Durbin-Watson test

### Homoscedasticity:

- Ensure constant variance in residuals
- Use Breusch-Pagan test

### Normality:

- Ensure residuals follow normal distribution
- Use QQ plot and Shapiro-Wilk test

### Handle assumption violations:

Use polynomial regression for non-linearity, remove correlated predictors, log-transform for heteroscedasticity, and apply robust regression for non-normal residuals.

# Checking Regression Assumptions

```
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.stats.diagnostic import het_breuschpagan
from scipy.stats import shapiro

# Linearity check
sns.residplot(x=y_pred, y=residuals, lowess=True)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.show()

# Normality check (QQ plot)
sm.qqplot(residuals, line='45')
plt.title('QQ Plot of Residuals')
plt.show()

# Shapiro-Wilk test
shapiro_test = shapiro(residuals)
print(f'Shapiro-Wilk Test p-value: {shapiro_test.pvalue}')

# Homoscedasticity test
_, pval, _, _ = het_breuschpagan(residuals, X)
print(f'Breusch-Pagan Test p-value: {pval}')
```

# Implementing Logistic Regression

---

# Logistic Regression with Statsmodels

Modelling Binary Outcomes with Statistical Insight



- Model binary outcomes based on predictors
- Provide detailed statistical summaries
- Support hypothesis testing and confidence intervals
- Interpret logistic regression results effectively
- Predict Titanic survival using passenger features

# Dataset Overview

## Features Used to Predict House Price

Variable	Description
Survived	Binary target variable (1 = Survived, 0 = Did Not Survive)
Pclass	Passenger class (1st, 2nd, 3rd)
Sex	Gender of passenger
Age	Age of passenger
SibSp	Number of siblings/spouses aboard
Parch	Number of parents/children aboard
Fare	Ticket fare
Embarked	Port of embarkation (C, Q, S)

- Build a linear regression model to classify whether a passenger survived or not

# Importing Libraries and Loading Data

## Preparing Titanic Dataset for Logistic Regression

### Import Necessary Libraries:

- Use pandas, NumPy and Statsmodels for data handling and modelling

### Import Preprocessing Tools:

- Use OneHotEncoder for categorical encoding

### Load the Dataset:

- Read Titanic data from CSV into pandas DataFrame

```
import pandas as pd  
import numpy as np  
import statsmodels.api as sm  
from sklearn.preprocessing import OneHotEncoder
```

### # Load dataset

```
data = pd.read_csv('titanic.csv')
```

# Handling Missing Values and Data Preprocessing

## Preparing Clean and Structured Input for the Model

### # Fill missing values

```
data['Age'].fillna(data['Age'].median(), inplace=True)
```

### # Drop irrelevant columns

```
data.drop(columns=['PassengerId', 'Name', 'Ticket', 'Cabin'],  
          inplace=True)
```

### # One-hot encode categorical variables

```
data = pd.get_dummies(data,  
                      columns=['Sex', 'Embarked'],  
                      drop_first=True)
```

# Selecting Features and Building the Model

## Preparing and Training a Logistic Regression Model

### # Select features and target

```
X = data[['Pclass', 'Age', 'SibSp', 'Parch', 'Fare', 'Sex_male', 'Embarked_Q', 'Embarked_S']]  
y = data['Survived']
```

### # Add constant term

```
X = sm.add_constant(X)
```

### # Build the model

```
model = sm.Logit(y, X)
```

### # Fit the model

```
results = model.fit()
```

# Model Summary and Interpretation

## Evaluating Logistic Regression Results

- **Log-likelihood:** Measure model fit
- **Pseudo R-squared:** Use as an alternative to R-squared
- **Coefficients and odds ratios:** Interpret effect of each predictor
- **P-values:** Determine predictor significance
- **Interpretation Example:**  
Odds ratio > 1 increases survival probability  
Odds ratio < 1 decreases survival probability

### # Code Implementation

```
print(results.summary())
odds_ratios = np.exp(results.params)
print(odds_ratios)
```

# Model Evaluation and Assumption Checks

## Assessing Predictions and Ensuring Model Validity

### Make predictions

```
y_pred_prob =  
results.predict(X)  
  
y_pred = [1 if prob >  
0.5 else 0 for prob in  
y_pred_prob]
```

### Calculate metrics

```
confusion_matrix(y,  
y_pred)  
  
accuracy_score(y,  
y_pred)
```

### Check assumptions

- Linearity of log-odds
- Independence of observations
- No perfect multicollinearity
- Large sample size

### Handle violations

- Use polynomial terms for non-linearity
- Remove correlated predictors
- Adjust for class imbalance

# Q & A

**Thank you**